

JeevanSetu

Finds the conflicts between a person's will and their bank, insurance, mutual fund and NPS nominations — while they are still alive to fix them.

FORMAT	Four-minute live demo
RULE SET	7 sourced rules, version 1.0.0-candidate
DEMO DATA	11 synthetic documents, 4 personas
TESTS	136 backend, 33 frontend
DEPLOYED	Render, behind a shared credential

READ THIS FIRST

Every document, person, institution and identifier in this project is fictional. Nothing here is legal advice, nothing is connected to a real bank, insurer, asset management company or the NPS, and no rule in it decides who legally inherits anything. The tool reports that two documents disagree and quotes both. That is all.

1. The one-line version

Every year, nominee-versus-will mismatches turn routine claims into family disputes and delayed settlements. We built the tool that catches them while someone is still alive to fix them.

2. The problem, as one concrete story

A person opens a bank account in 2023 and names their spouse as nominee. In 2026 they write a will leaving that same account to their daughter. Both documents are valid. Neither references the other. The bank has never seen the will, and the will was never sent to the bank.

Nobody notices until after a death, when the bank settles on its own record and the family discovers the will says something different. That is not a paperwork annoyance. That is the start of a dispute.

This happens because every institution in India — banks, insurers, mutual funds, the NPS — independently collects nominee data, and none of them cross-check against a person's will or against each other. The gap is structural, not a one-time inefficiency.

3. What we built

- Upload a will and the nomination forms. The tool:
- 1. Classifies each document and extracts who is named for each asset, under which mechanism, with the exact source text and a locator.
 - 2. Cross-references every document for the same person against a small, sourced rule set.
 - 3. Reports each conflict with both sources quoted, a plain-language explanation, a suggested fix, and the limits of what it can conclude.
 - 4. Produces a Continuity Readiness Score, so there is a reason to come back after each fix.

4. The demo, beat by beat

BEAT	WHAT THE JUDGE SEES
Score	75 / 100 for the conflict persona, with the arithmetic shown on screen: 100 minus 25 for one high-priority conflict.
The finding	One conflict card: <code>BANK-WILL-001</code> on account <code>SB-DEM0-4821</code> . Source A is the bank nomination naming Kavya. Source B is will clause 3.1 naming Anika. Both quoted verbatim from the uploaded PDFs, with the rule ID and the scope note.
Fix it live	Change the nominee in the vault, press re-run: the score goes to 100 and the conflict clears.
Undo it	Press undo: the original extracted value comes back and the score returns to 75 . The edit stays in the log.
Contrast	A second persona whose documents agree scores 100 with no findings.

The whole demo runs from a fresh clone in about 40 seconds with no API key, using a clearly labelled pre-verified extraction snapshot.

Two more personas, for the questions that follow

These exist to show that the tool distinguishes three different kinds of answer rather than labelling everything a conflict.

PERSONA	WHAT IT DEMONSTRATES	SCORE
Priya Demo-Iyer validity warning	Her NPS nomination predates her confirmed marriage date, so <code>NPS-VALIDITY-001</code> questions that one document's standing rather than reporting a disagreement between two. The will-mismatch rule is deliberately suppressed here instead of stacked on top, which is the rule ordering made visible.	75
Meera Demo-Rao manual review required	Her nomination and will name different people for the same account, which would normally be <code>BANK-WILL-001</code> . It is not, because the account is held jointly and that rule covers individually held accounts. The tool reports <code>GAP-JOINT-HOLDING</code> and declines to judge.	97

Say the 97 out loud before a judge spots it: leaving rule scope *raises* the score, because a gap deducts 3 where a conflict deducts 25. A gap is not a clean bill of health, and the card says manual review required.

5. The credibility answers

"Is the AI deciding what counts as a conflict?"

No. AI extracts fields and rewords a finding after the fact. Every rule decision is deterministic Python in one file, `detection.py` , which imports no model client — and a test asserts both that its source contains no AI reference and that constructing a model client during detection fails. That is why we can show you the rule that fired and the exact text it fired on.

"Are these real legal rules?"

Seven rules, each sourced to a Supreme Court judgment, the Insurance Act, a SEBI circular or the PFRDA regulations, all written down in `RULES.md` with the predicate, the severity, the exclusions and worked examples. We re-verified every source before the demo. We deliberately encoded nothing we could not source: SEBI published a March 2026 consultation paper proposing new nomination norms, and because it is still a draft, no rule is based on it. Where the law is unsettled, the tool says "manual review required" instead of guessing.

"Does it tell people who legally inherits?"

Never. That is the discipline that makes it usable. It says these two documents disagree, here is the text, here is who to ask. Every result carries a not-legal-advice disclaimer, and model wording is rejected outright if it contains a legal conclusion — phrases like "legally owns", "overrides the will" or "is the legal heir" fail a deterministic check before the text ever reaches the screen.

"What happens when the API fails?"

The upload is kept, the document is marked failed, and it stays retryable. If AI wording fails, the finding still appears with its deterministic summary and both sources. For the demo we ship a pre-verified fallback snapshot, and the interface labels it as "pre-verified fallback data, not a live AI extraction" rather than passing it off as live.

"Can you undo what you just changed?"

Yes. Every simulated fix is logged with its previous and new value, and undo unwinds edits one at a time. Re-extraction refuses to discard your edits unless explicitly forced. A fix only ever changes the extracted synthetic record — never an uploaded document, never an institution's record.

6. Why this is a company, not a feature

- **The wedge:** a sharp, free-to-try tool that finds real conflicts. The value is immediate and needs no institutional partnership.
- **It compounds:** fixing one conflict pulls the user into adding the next account, and life events — marriage, a child, a new policy — invalidate old nominations. Re-use is built into the problem, not bolted on.
- **Two expansion paths that need no integration:** paid deeper reviews and family access on the consumer side; an embeddable pre-claim check for banks, insurers and wealth managers who also lose money to disputed claims. We never need write access to anyone's account — only documents the user chooses to share.
- **The second act:** settlement workflows become viable later precisely because the hard data problem — incomplete, conflicting nominee data — was already solved earlier in the user's life.

7. What we are not claiming

- Not connected to any bank, insurer, asset management company or the NPS.
- Not a legal opinion, and not exhaustive of Indian succession law.
- Not tested on real documents. Every document in the demo is synthetic and marked as such, and the upload path requires an explicit synthetic attestation.
- Joint holdings are recognised and deliberately excluded from the rules: a joint asset is reported as "manual review required", not as a conflict. A document that does not state its holding pattern is still checked, and that residual risk is written down.
- Four of the seven rules are not exercised by any seeded persona and are covered only by unit tests. The demo shows one rule firing, one suppressed, one gap, and one clean result.
- The hosted instance has no user accounts, no persistent storage, and one shared password. It sleeps after 15 minutes idle.
- The per-institution settlement checklist is labelled **Concept / Future Feature** in the API payload and on its own page. It is an illustration, not an integration, and it touches neither the rule engine nor the score.

8. If asked what is next

1. Life-event reminders, since a marriage silently invalidates certain NPS nominations.
2. A reviewed expansion of the rule set, one sourced rule at a time, starting with the asset types currently excluded.
3. A hosted demo behind real access control.
4. Postgres and a migration path, replacing the single-file SQLite the prototype resets on a schema change.

9. The numbers, if you are asked

CLAIM	FIGURE
Sourced deterministic rules	7
Labelled out-of-scope markers (never presented as rules)	7
Fields extracted per named person, each with source text	18
Synthetic demo documents	11
Demo personas, one per possible outcome	4
Backend tests	136
Frontend tests	33
Fresh clone to a working demo	~40 s
AI calls involved in any conflict decision	0

JeevanSetu is a hackathon prototype built on synthetic documents. Not legal advice. The rule set is a limited, India-scoped first pass; a licensed legal professional and the relevant institution should review any real document.